

Artificial Intelligence PBL-1

Classifier for MNIST with LR, and SVM

Team 4

Given Problem

▶ Problem : making an automated zipcode recognition system

▶ Main objective:

Find the best model out of the following settings:

Classifiers: Logistic Regression and SVM

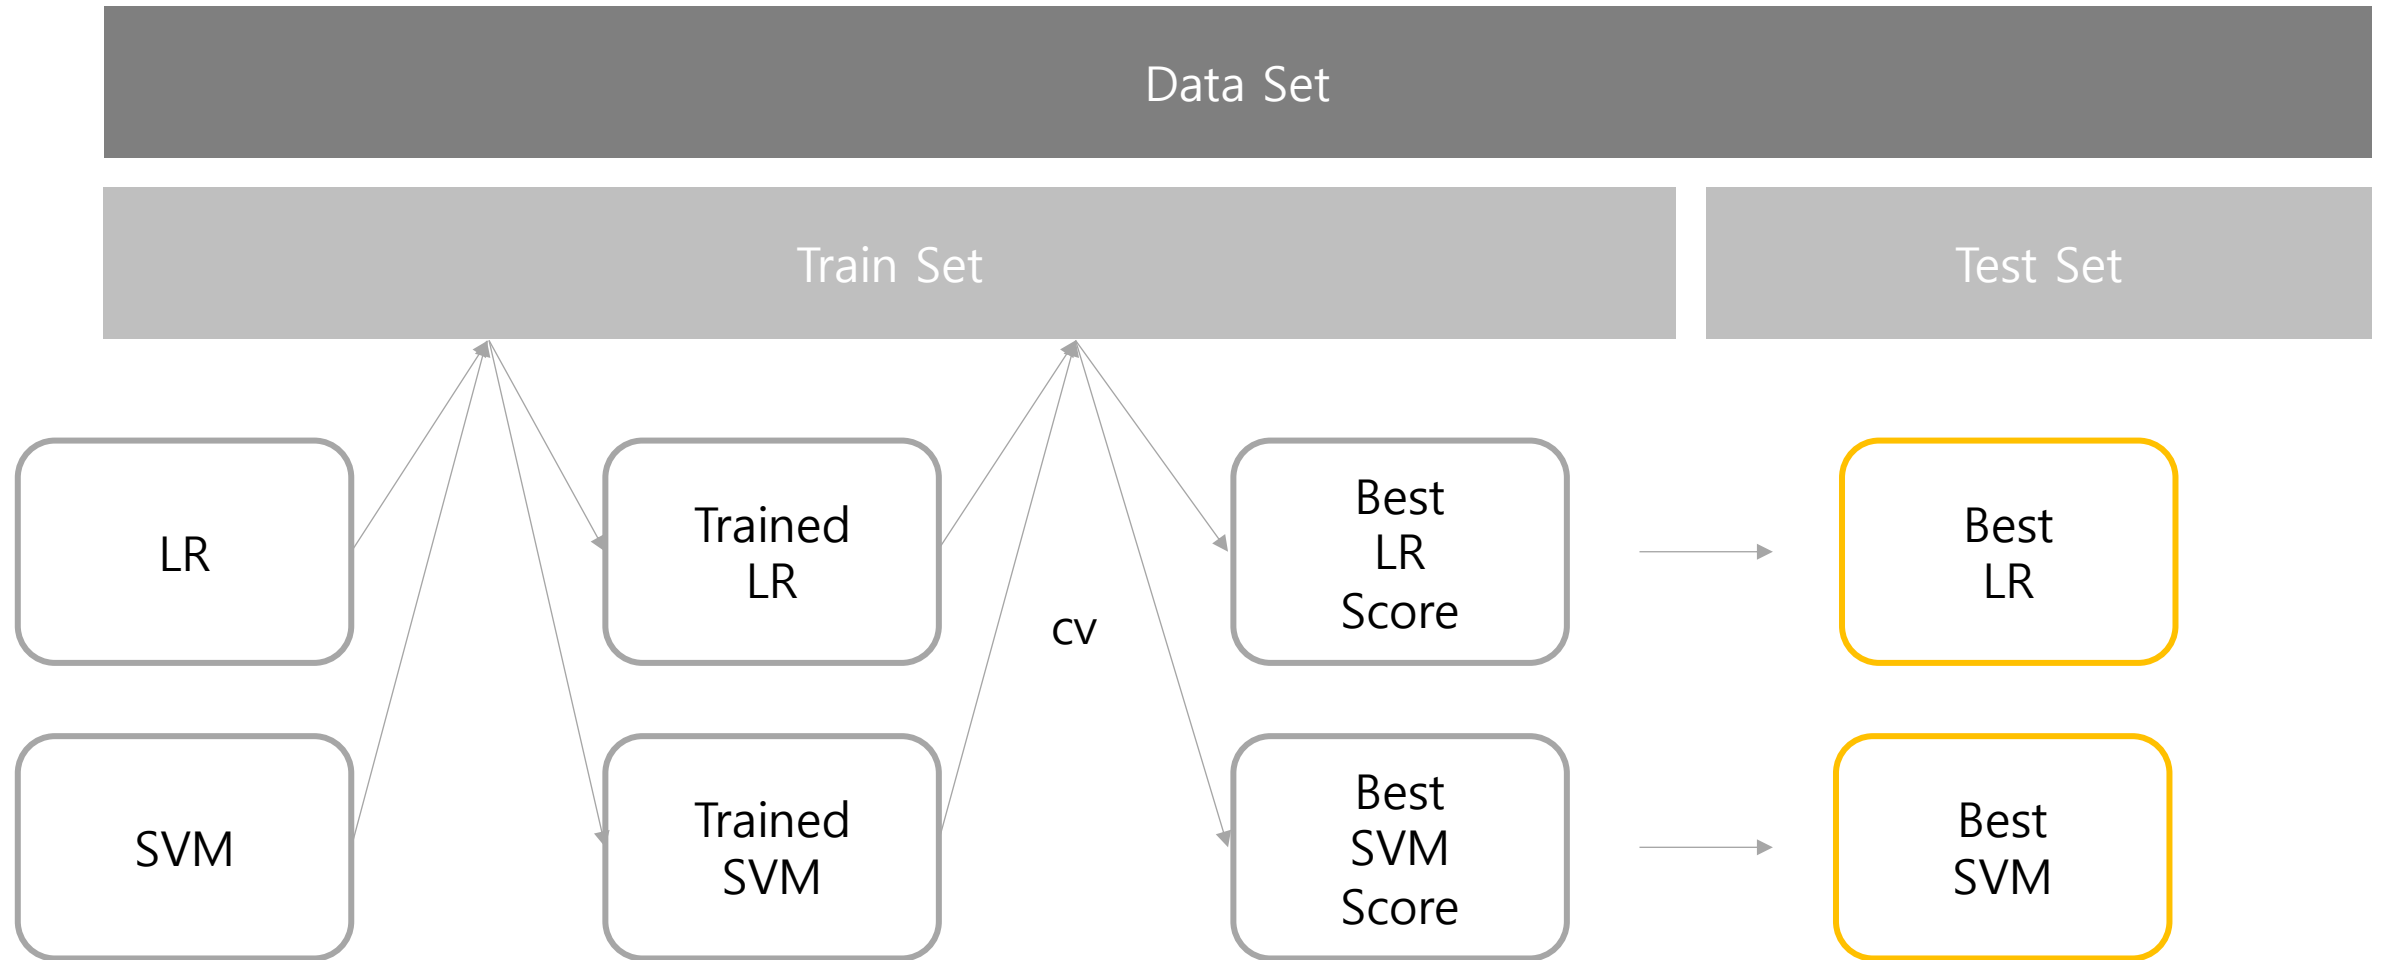
Data: 70k number(0~9) images of MNIST

Criterion:

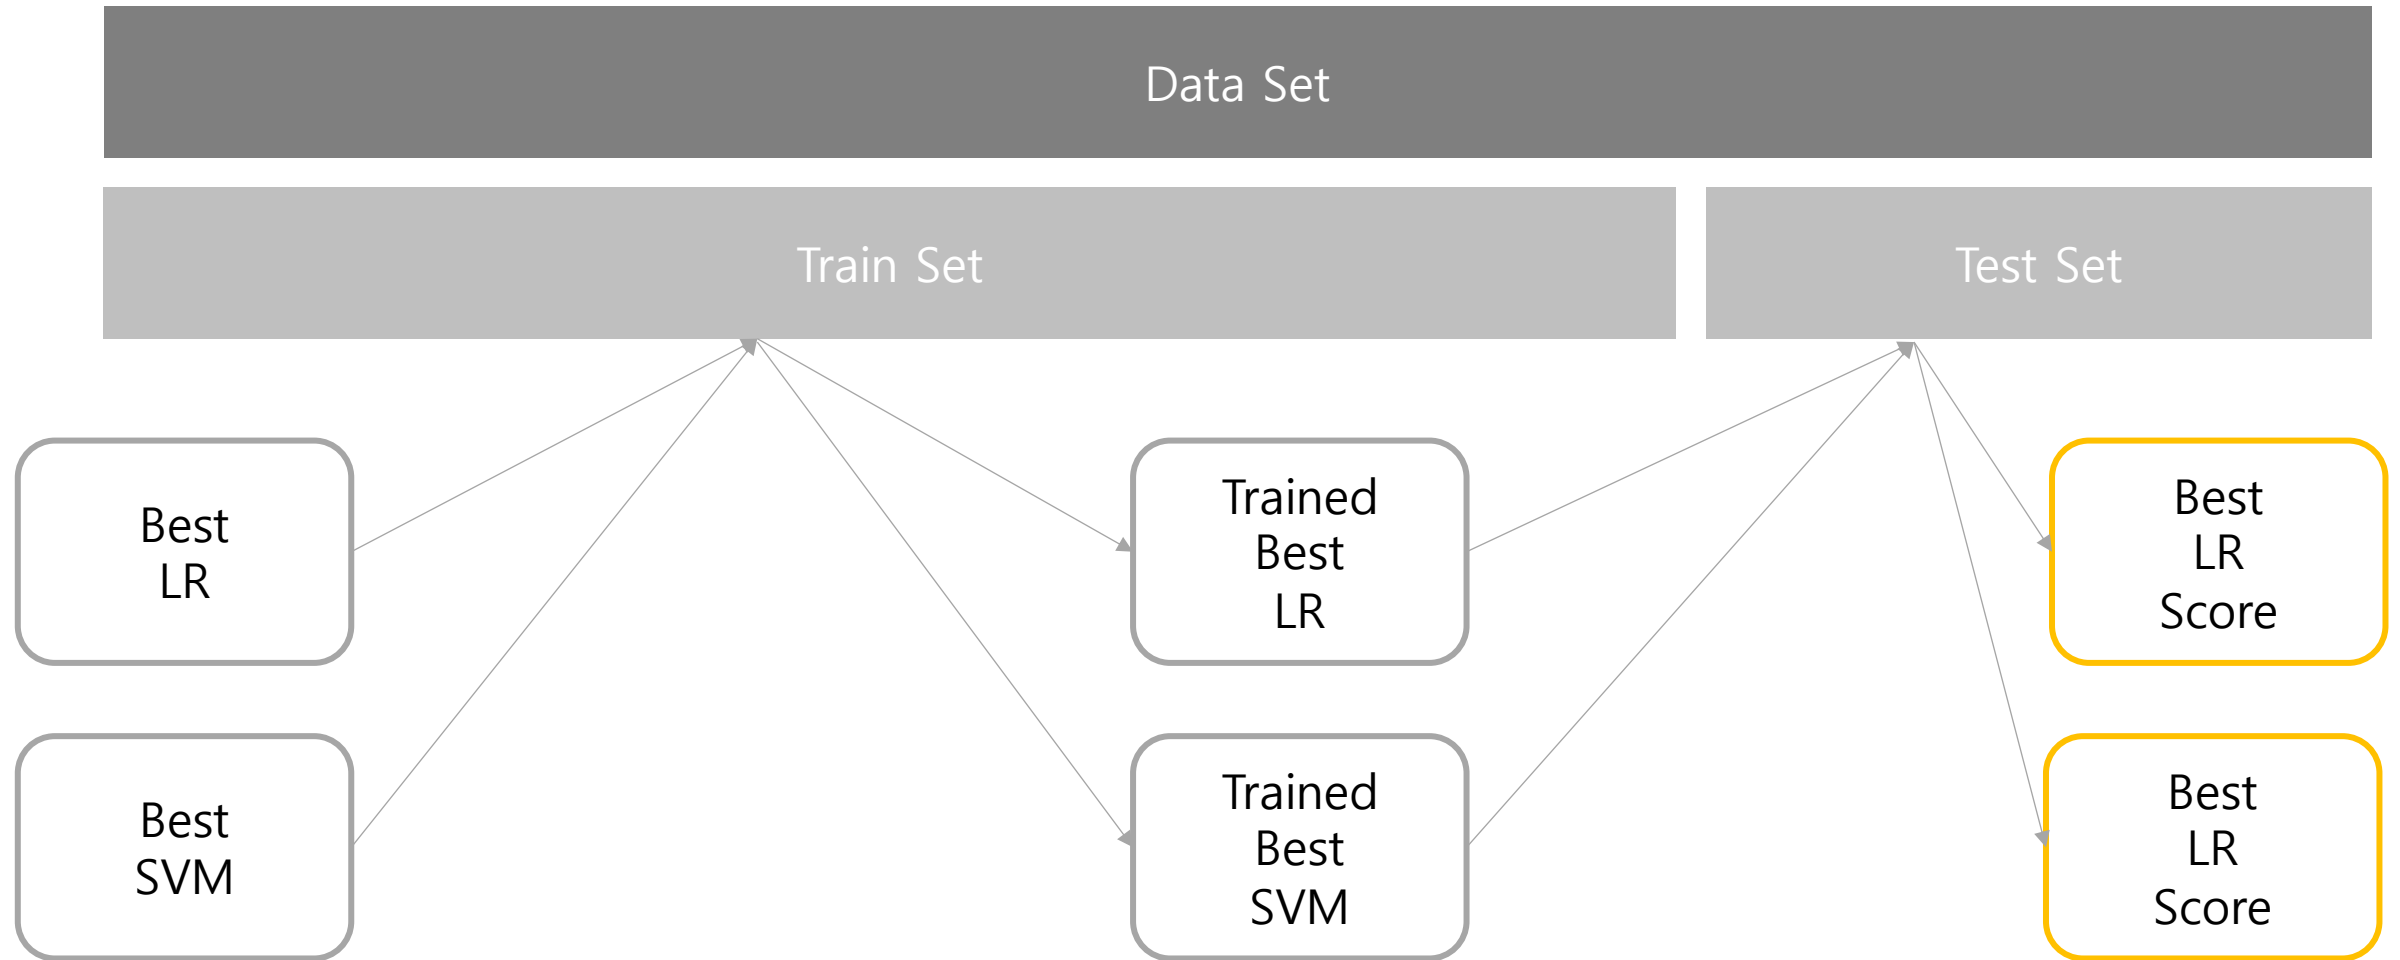
Model selection: f1-score(micro)

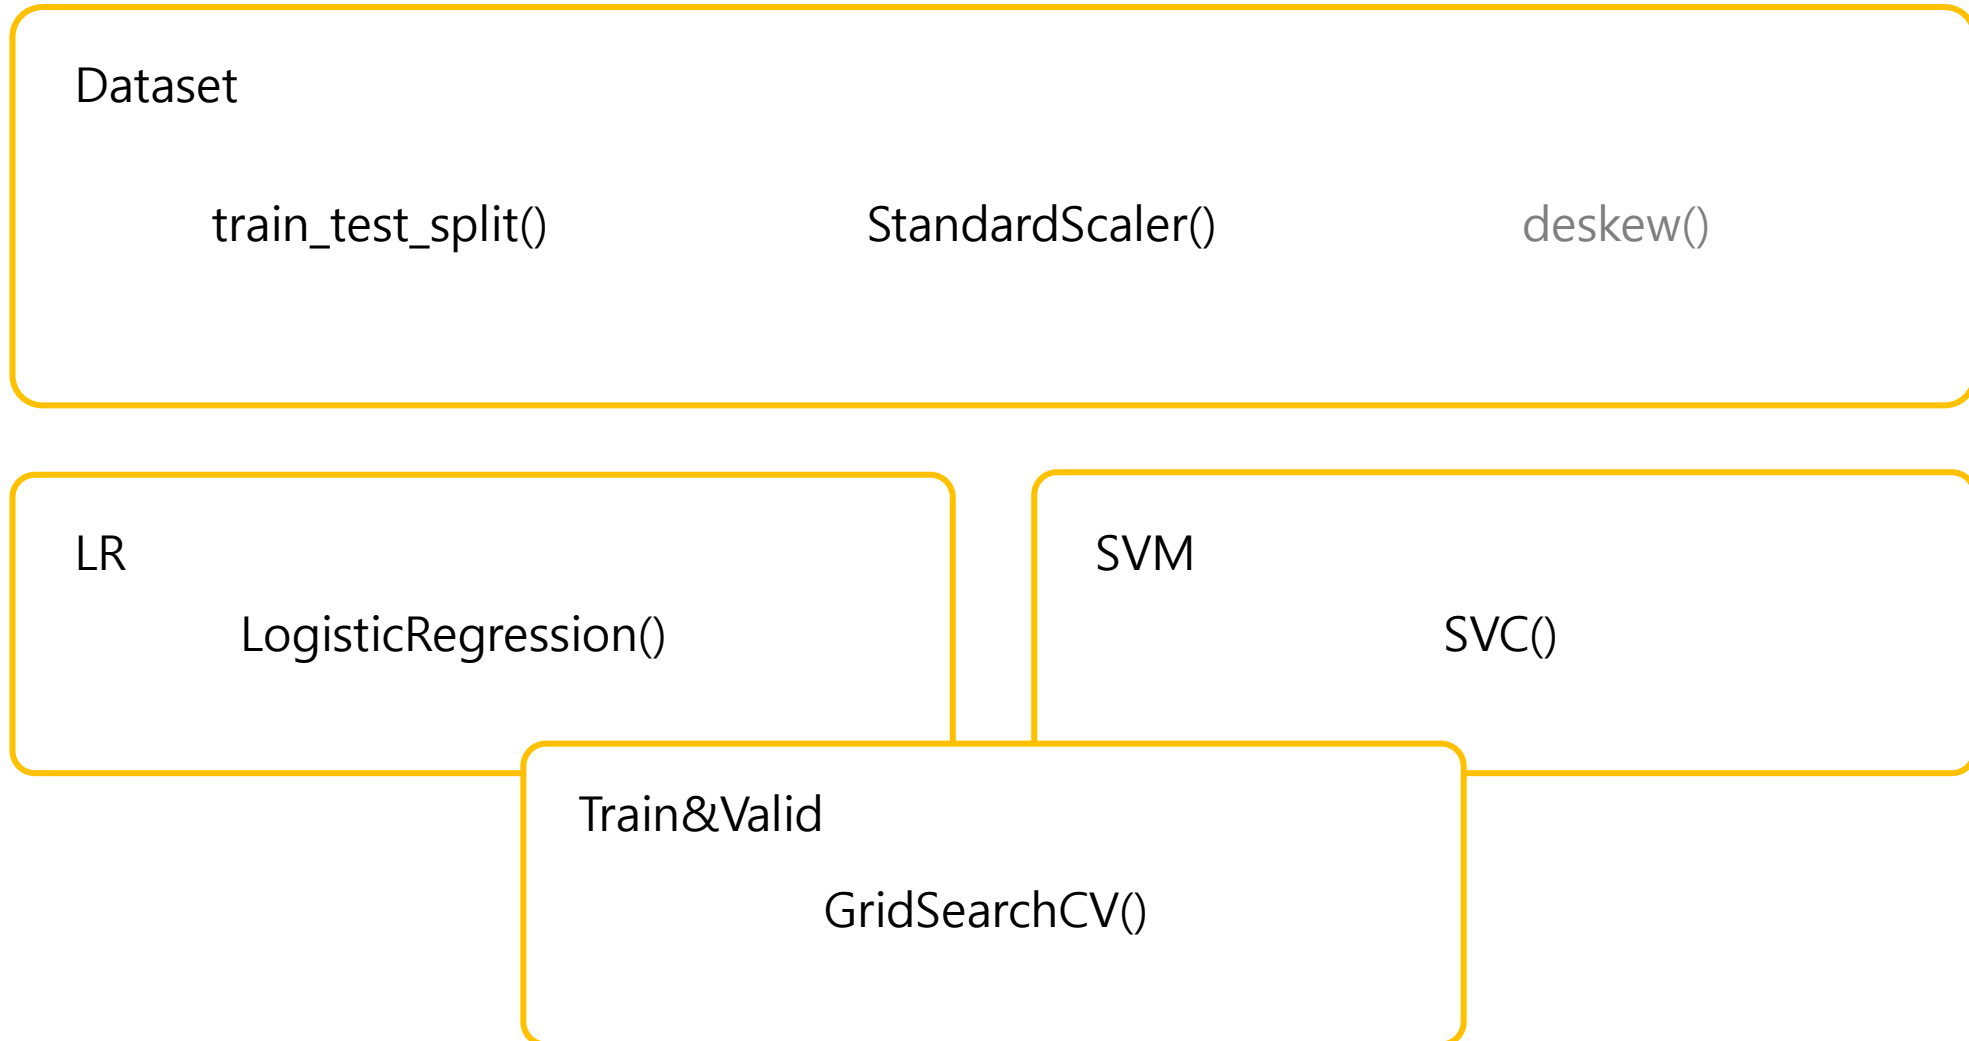
Test performance: f1-score(micro) of the final model on the test images

Solution



Solution






deskew()

y-critical system
distinction can be
tested. The mission
behaviour while the
y controller when
more, the aims of
mission controller
are – this will also
lead into an unsafe
state with avoiding
unsafe states that


deskew

y-critical system
distinction can be
tested. The mission
behaviour while the
y controller when
more, the aims of
mission controller
are – this will also
lead into an unsafe
state with avoiding
unsafe states that

do not(acc : 0.944375)




do(acc : 0.952)



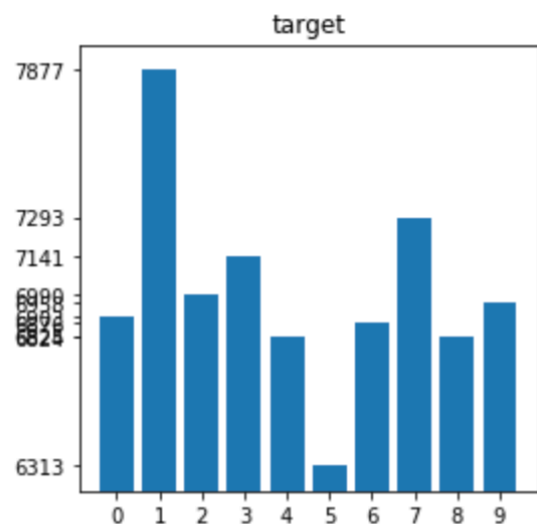
MNIST DATA

- 28x28 pixel
- pixel value range [0, 255]
- 70000 samples
- 786 features
- many pixel values are zero
- large '1's, small '5's
- no outlier, missing value



Dataset

from sklearn.model_selection import train_test_split

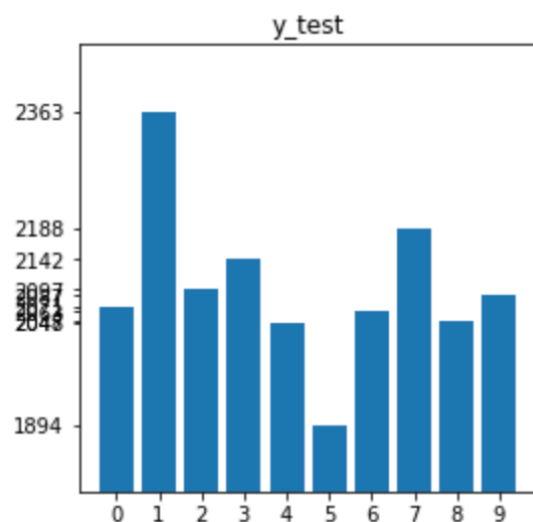
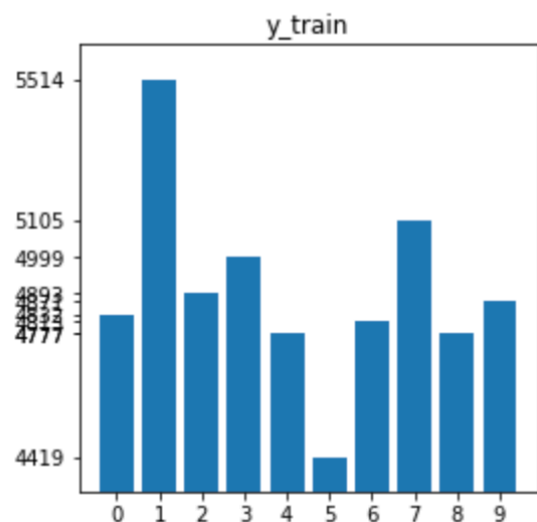


X_train : (49000, 784)

X_test : (21000, 784)

y_train : (49000,)

y_test : (21000,)



```
mnist = fetch_openml("mnist_784")
X = mnist.data
y = mnist.target

X_train, X_test, y_train, y_test
= train_test_split(X, y, test_size=0.3,
                    random_state=123, stratify=y)
```



```
from sklearn.preprocessing import StandardScaler
```



```
scaler = StandardScaler()  
scaler.fit(X_train)  
X_train = scaler.transform(X_train)  
X_test = scaler.transform(X_test)
```

Scaler:

StandardScaler() # mean: 0, std: 1

MinMaxScaler() # max: 1, min: 0

RobustScaler() # median: 0, IQR: 1

Standard score : $\frac{X - \mu}{\sigma}$

Min-Max Feature scaling : $\frac{X - X_{min}}{X_{max} - X_{min}}$

```
from sklearn.model_selection import GridSearchCV
```

```
scoring = f1_micro  
cv = 10
```

$$F_1 = \left(\frac{2}{\text{recall}^{-1} + \text{precision}^{-1}} \right) = 2 \cdot \left(\frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}} \right)$$

- ▶ **f1_score – macro** : Calculate metrics for each label, and find their unweighted mean.
This does not take label imbalance into account.
 - **micro** : Calculate metrics globally by counting the total true positives, false negatives and false positives.
- ▶ **cv** : Determines the cross-validation splitting strategy.

1	2	3	4	5	6	7	8	9	10
---	---	---	---	---	---	---	---	---	----

Logistic Regression



```
parameters = {'solver':['saga'], 'penalty':['l1', 'l2'], 'C':[100.0, 10.0, 1.0, 0.1, 0.01]}

lc = LogisticRegression(multi_class='ovr')
clf = GridSearchCV(estimator=lc, param_grid=parameters, scoring='f1_micro', cv=10, n_jobs=-1)
clf.fit(X_train, y_train)
```

Support Vector Machine



```
parameters = {'kernel':['linear'], 'C':[100.0, 10.0, 1.0, 0.1, 0.01, 0.001]}

svc = SVC()
clf = GridSearchCV(estimator=svc, param_grid=parameters, scoring='f1_micro', cv=10, n_jobs=-1)
clf.fit(X_train, y_train)
```

```
from sklearn.linear_model import LogisticRegression
```

LogisticRegression() : Logistic Regression classifier.

- multi_class : ovr(one-vs-rest)
- penalty : l1, l2

$$\|x\|_p := \left(\sum_{i=1}^n |x_i|^p \right)^{1/p} \quad \begin{array}{l} \text{if } p = 1, \text{ l1} \\ \text{if } p = 2, \text{ l2} \end{array}$$

▶ best parameter :

```
penalty = l2  
C = 100.0  
solver = saga
```

- C : 100.0, 10.0, 1.0, 0.1, 0.01
- solver : saga

$$w^{k+1} = w^k - \alpha (f'_j(w^k) - f'_j(\theta_j^k)) + \frac{1}{n} \sum_{i=1}^n f'_i(\theta_i^k)$$


▶ best score :

```
f1_micro = 0.892
```

Validation

- Logistic Regression

C \ f1		mean_test_score	mean_fit_time(min)
100.0	l1	0.89189796	14
	l2	0.892	10
10.0	l1	0.89195918	15
	l2	0.89191837	10
1.0	l1	0.89193878	15
	l2	0.89197959	10
0.1	l1	0.89159184	16
	l2	0.89193878	10
0.01	l1	0.88708163	14
	l2	0.8919837	11



Intel Xeon, RAM 13GB

```
from sklearn.svm import SVC
```

SVC() : C-Support Vector Classification

- kernel : linear
- C : 100.0, 10.0, 1.0, 0.1, 0.01, 0.001

▶ best parameter :

C = 0.01


▶ best score :

f1_micro = 0.94155102

Validation

- SVM

C \ f1	mean_test_score	mean_fit_time(min)
100.0	0.91259184	215
10.0	0.91573469	23
1.0	0.91922449	10
0.1	0.93163265	9
0.01	0.94144102	9
0.001	0.93922449	11



AMD 2700x, RAM 32GB

LR

multi_class = ovr

penalty = l2

C = 100.0

solver = saga

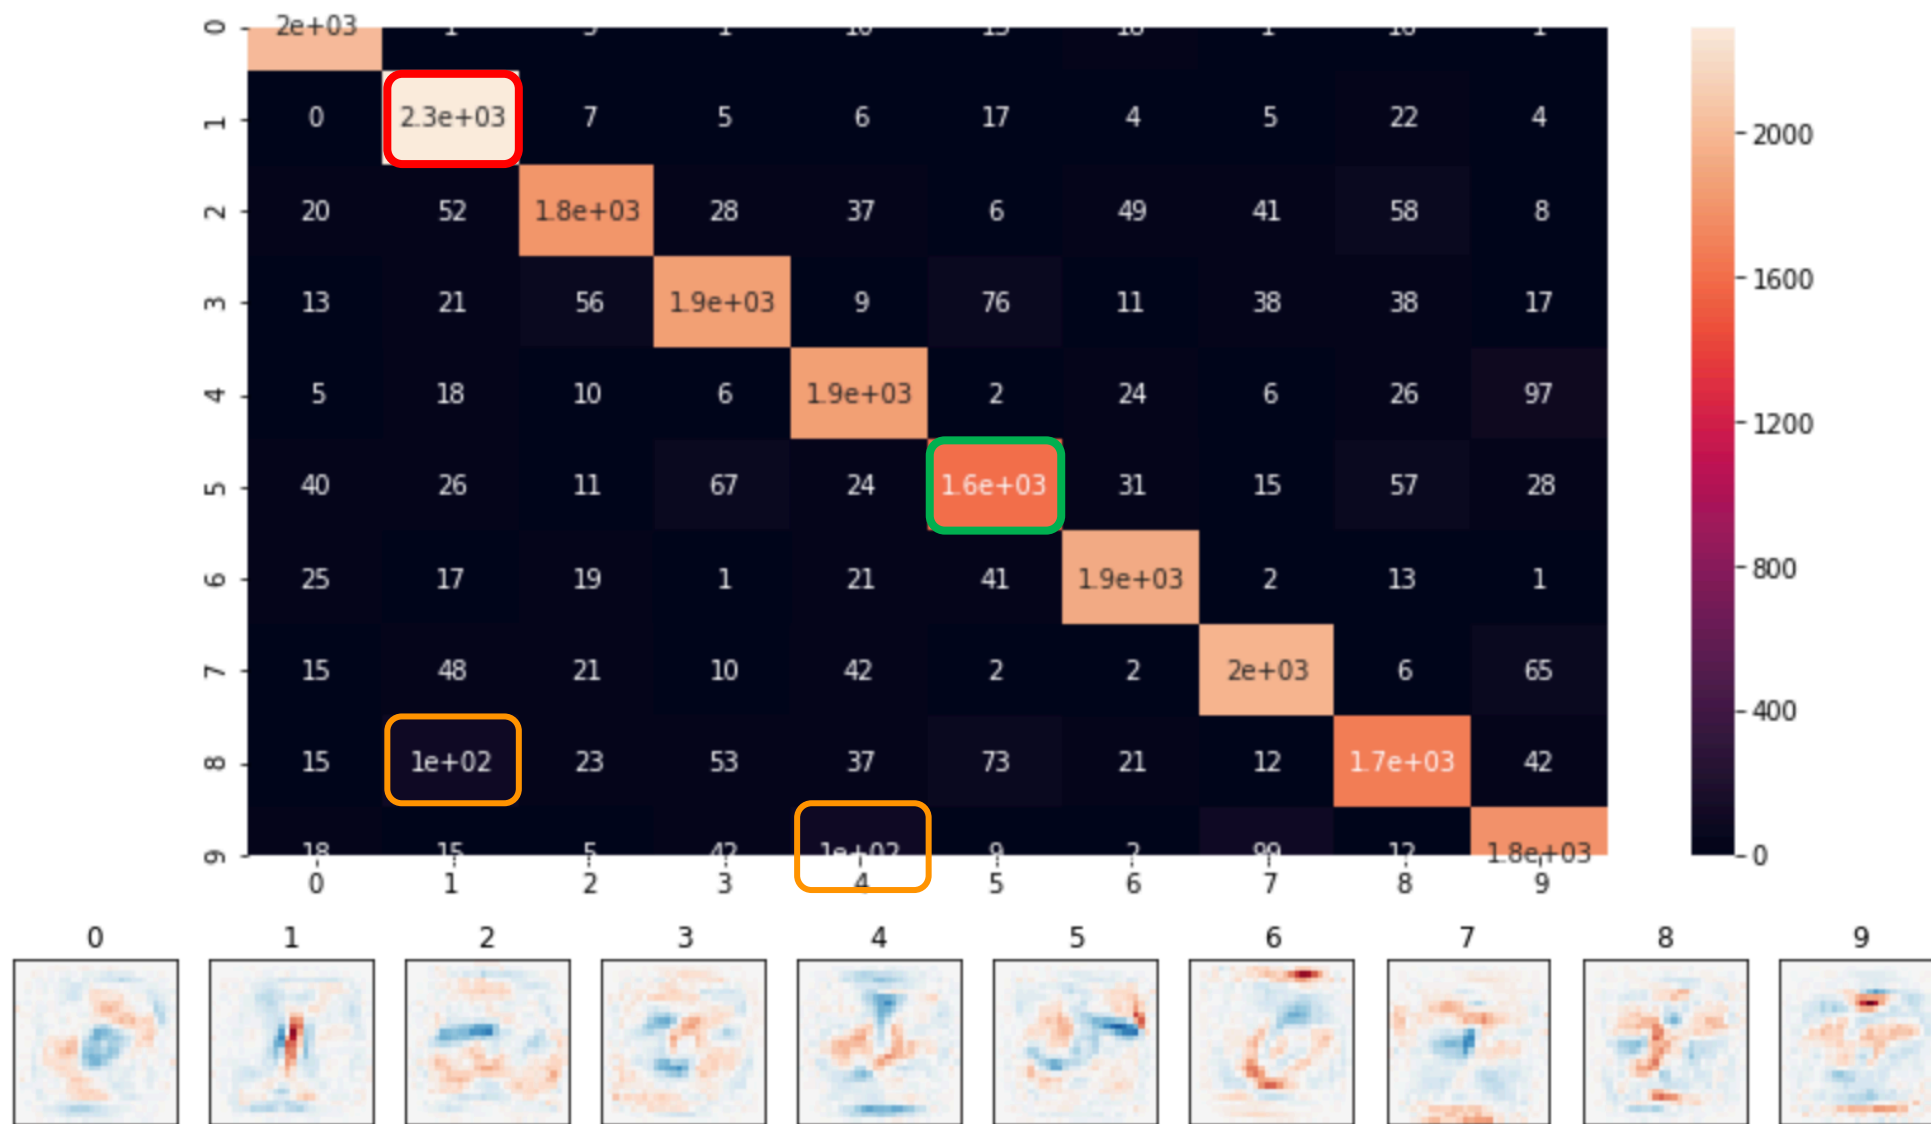
SV
M

kernel = linear

C = 0.01

Test

- LR




- LR

	precision	recall	f1-score	support
0	0.93	0.97	0.95	2071
1	0.88	0.97	0.93	2363
2	0.92	0.86	0.89	2097
3	0.90	0.87	0.88	2142
4	0.87	0.90	0.89	2047
5	0.87	0.84	0.86	1894
6	0.92	0.93	0.93	2063
7	0.90	0.90	0.90	2188
8	0.87	0.82	0.84	2048
9	0.87	0.86	0.86	2087
micro avg	0.89	0.89	0.89	21000
macro avg	0.89	0.89	0.89	21000
weighted avg	0.89	0.89	0.89	21000

Test

- SVM



- SVM

	precision	recall	f1-score	support
0	0.96	0.98	0.97	2071
1	0.96	0.98	0.97	2363
2	0.93	0.94	0.93	2097
3	0.92	0.92	0.92	2142
4	0.93	0.95	0.94	2047
5	0.90	0.91	0.91	1894
6	0.97	0.96	0.96	2063
7	0.96	0.93	0.95	2188
8	0.94	0.91	0.92	2048
9	0.93	0.93	0.93	2087
micro avg	0.94	0.94	0.94	21000
macro avg	0.94	0.94	0.94	21000
weighted avg	0.94	0.94	0.94	21000

- micro

	Accuracy	Precision	Recall	F1-score
LR	0.893619	0.893619	0.893619	0.893619
SVM	0.941905	0.941905	0.941905	0.941905

- macro

	Accuracy	Precision	Recall	F1-score
LR	0.893524	0.89332	0.892042	0.892141
SVM	0.941905	0.941301	0.941153	0.941147

$$ACC = \frac{TP + TN}{ALL}$$

$$PRE_{micro} = \frac{TP_1 + \dots + TP_k}{TP_1 + \dots + TP_k + FP_1 + \dots + FP_k}$$

$$REC_{micro} = \frac{TP_1 + \dots + TP_k}{TP_1 + \dots + TP_k + FN_1 + \dots + FN_k}$$

$$F1 = 2 \cdot \frac{PRE \times REC}{PRE + REC}$$

Question and Answer

A solid yellow vertical bar is positioned on the left side of the slide.

Thank you